

Paper 2 Revision Memo

Completed updates

The revised manuscript implements the recommended external-submission revisions: a formal EIT pipeline figure, expanded literature integration, explicit limitation clarity, protocol-level pseudocode, schema examples, an audit traceability matrix, and test vectors.

Visual documentation

Figure 1 now presents the pipeline Consent Scope -> LIL -> TAR -> PDEV -> Audit Traceability. The figure is embedded in the manuscript and included as a separate PNG file.

Literature integration

The reference list now includes 27 sources across AI risk management, EU AI regulation, OECD principles, RFC conformance language, internal algorithmic auditing, model documentation, datasets documentation, socio-technical infrastructure, and AI ethics governance.

Limitation clarity

The manuscript explicitly states that EIT governs longitudinal interaction telemetry and does not address model training data, compute allocation, model weight custody, frontier capability thresholds, or supply-chain security.

Protocol-level determinism

Appendix A provides pseudocode. Appendix B provides schema examples and the package includes machine-readable JSON schema files. Appendix C maps rules to audit evidence and test vectors.